

Yousha Murhij, BMSTU, Coding report1

Understanding the original dataset:

The original one batch data is (10000 x 3072) matrix expressed in numpy array. The number of columns, (10000), indicates the number of sample data. As stated in the CIFAR-10/CIFAR-100 dataset, the row vector, (3072) represents an color image of 32x32 pixels.

Regenerating the results of the paper for CIFAR10 Dataset:

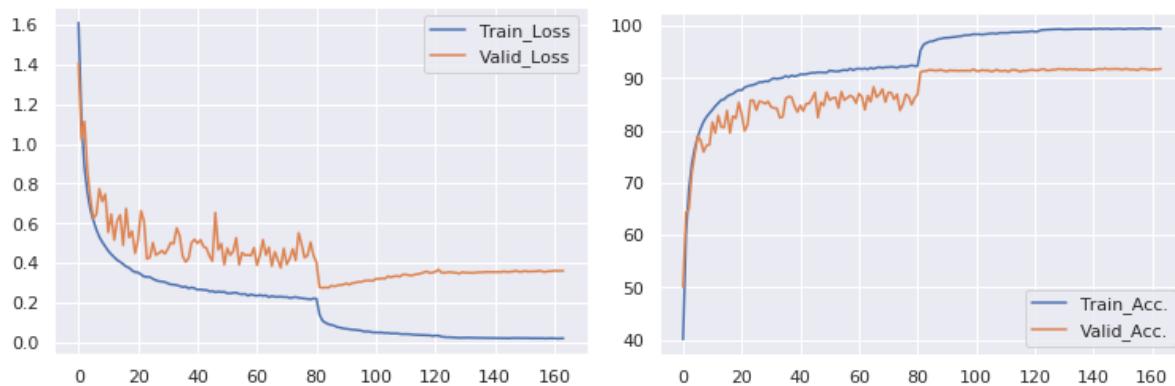
RESNET 20: 164 Epochs

Epoch: [164 | 164] LR: 0.001000

Processing (391/391) Data: 0.015s | Batch: 0.069s | Total: 0:00:27 | ETA: 0:00:01 | Loss: 0.0209 | top1: 99.3900 | top5: 99.9980

Processing (100/100) Data: 0.010s | Batch: 0.032s | Total: 0:00:03 | ETA: 0:00:01 | Loss: 0.3605 | top1: 91.7600 | top5: 99.7000

Best acc: tensor (91.8100, device='cuda:0')



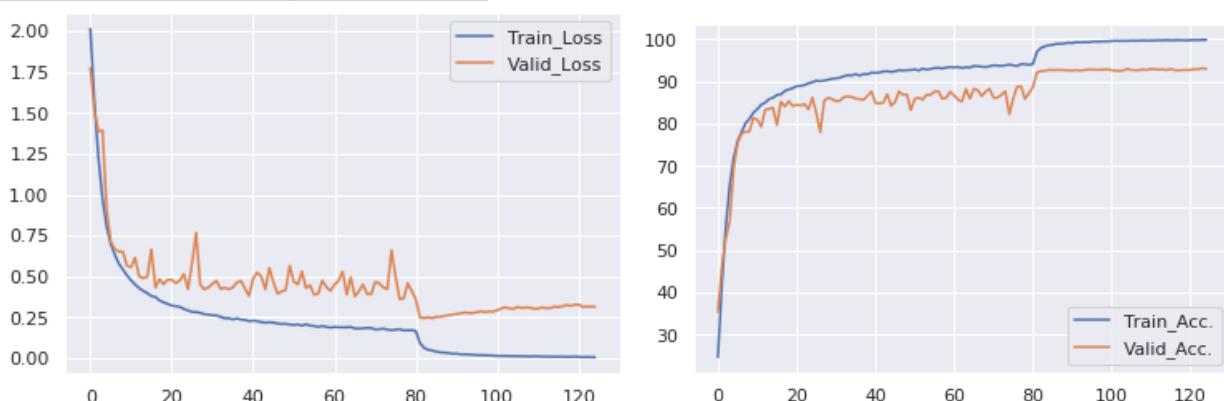
RESNET 56: 125 Epochs

Epoch: [125 | 125] LR: 0.001000

Processing (391/391) Data: 0.064s | Batch: 0.161s | Total: 0:01:03 | ETA: 0:00:01 | Loss: 0.0061 | top1: 99.8740 | top5: 100.0000

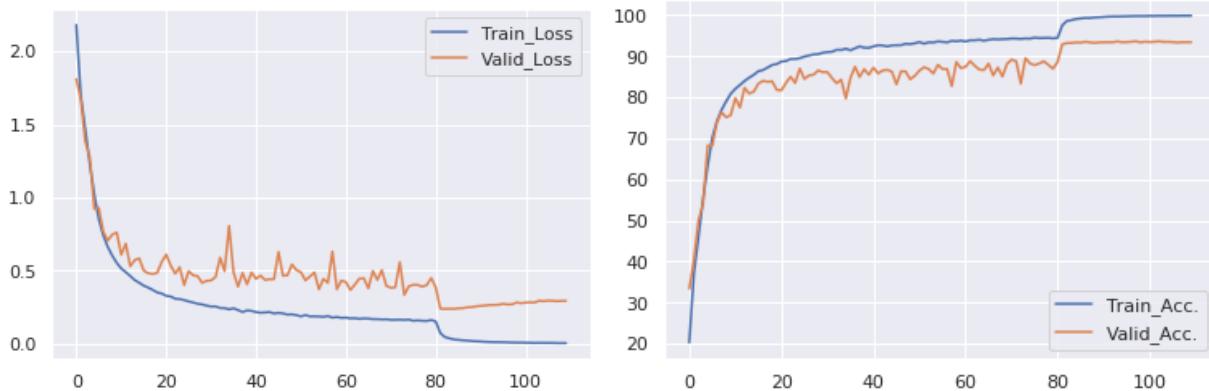
Processing (100/100) Data: 0.012s | Batch: 0.046s | Total: 0:00:04 | ETA: 0:00:01 | Loss: 0.3145 | top1: 93.0100 | top5: 99.7900

Best acc: tensor (93.1600, device='cuda:0')



RESNET 110: 110 Epochs

```
Epoch: [110 | 110] LR: 0.010000
Processing (391/391) Data: 0.005s | Batch: 0.314s | Total: 0:02:02 | ETA: 0:00:01 | Loss: 0.0059 | top1:
99.8540 | top5: 100.0000
Processing (100/100) Data: 0.031s | Batch: 0.083s | Total: 0:00:08 | ETA: 0:00:01 | Loss: 0.2947 | top1:
93.4000 | top5: 99.8300
Best acc: tensor (93.6100, device='cuda:0')
```



While Training my model I noticed that we may get different results when training our models with different random seed.

For CIFAR10:

In case of 20 layers and setting the number of epochs equal to 164 as in the base-code and without changing other hyper parameters, I got the same accuracy in both training and testing as the one indicated in the paper.

In case of 56 layers, the training takes more time than the previous one and for a number of epochs equal to 125, the best accuracy I got was 93.210% for the testing data and 99.5760% for the training data

In the case of 110 layers, it took certainly more and more time to converge and for 110 epochs without changing anything else I got 99.6040% accuracy for training data and 93.6100% for the testing data

So, I can conclude that this deep net needs a sufficient number of epochs to converge and the results that I got are good enough and almost similar to the ones in the original paper and the small difference is due to the different number of epochs.

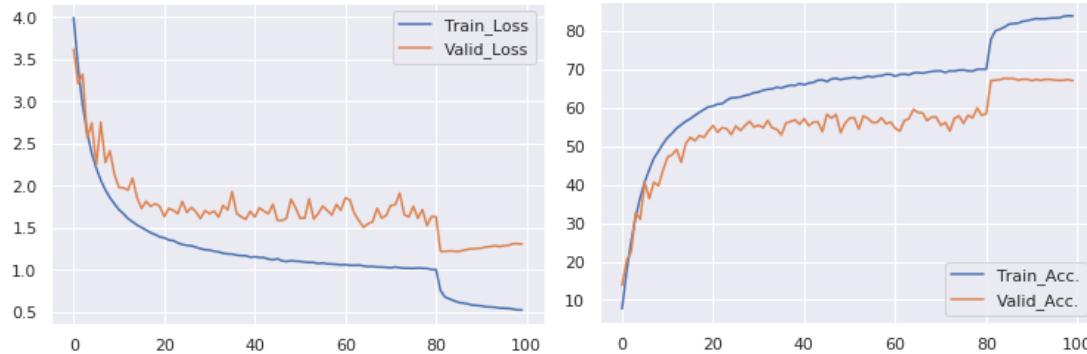
The most important notice that we should consider is that the needs about 80 epochs to converge and after that the weights, accuracy and the performance become stable and change slowly in case of additional epochs

Regenerating the results of the paper for CIFAR100 DataSet:

RESNET 20: 100 Epochs

```
Epoch: [100 | 100] LR: 0.010000
Processing (391/391) Data: 0.023s | Batch: 0.063s | Total: 0:00:24 | ETA: 0:00:01 | Loss: 0.5211 | top1:
83.8820 | top5: 97.8800
```

Processing (100/100) Data: 0.009s | Batch: 0.030s | Total: 0:00:03 | ETA: 0:00:01 | Loss: 1.3038 | top1: 67.0400 | top5: 90.3900
 Best acc: tensor (67.6700, device='cuda:0')



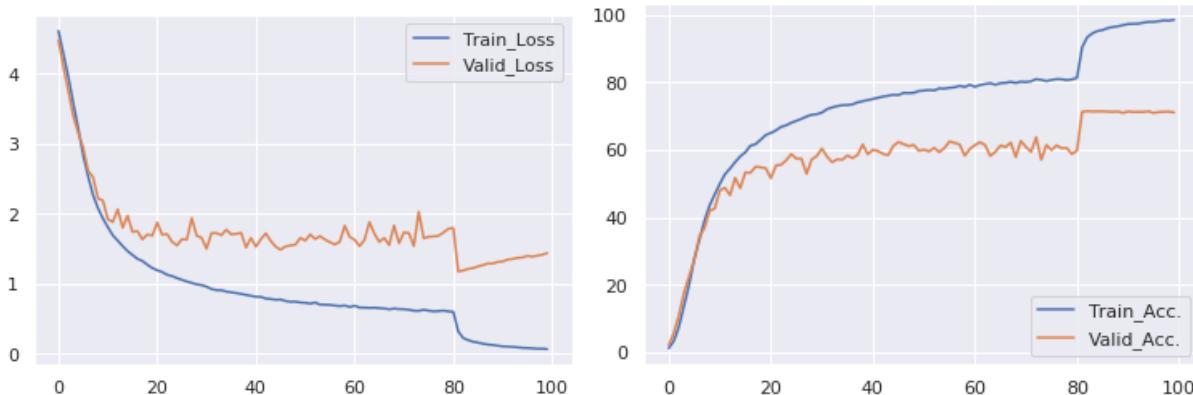
RESNET 56: 100 Epochs

Epoch: [100 | 100] LR: 0.010000
 Processing (391/391) Data: 0.077s | Batch: 0.162s | Total: 0:01:03 | ETA: 0:00:01 | Loss: 0.1473 | top1: 95.6540 | top5: 99.8180
 Processing (100/100) Data: 0.017s | Batch: 0.042s | Total: 0:00:04 | ETA: 0:00:01 | Loss: 1.4247 | top1: 70.0800 | top5: 91.2400
 Best acc: tensor (71., device='cuda:0')



RESNET 110: 100 Epochs

Epoch: [100 | 100] LR: 0.010000
 Processing (391/391) Data: 0.005s | Batch: 0.305s | Total: 0:01:59 | ETA: 0:00:01 | Loss: 0.0671 | top1: 98.4240 | top5: 99.9700
 Processing (100/100) Data: 0.030s | Batch: 0.083s | Total: 0:00:08 | ETA: 0:00:01 | Loss: 1.4395 | top1: 71.0600 | top5: 91.7900
 Best acc: tensor (71.4200, device='cuda:0')



For CIFAR100 we get similar results to the ones we got on CIFAR10 except for that the gap between the Training acc. And the Validation became bigger and this happens because we increased the number of classes to 100 which leads to a bigger error margin.

Optional task 1:

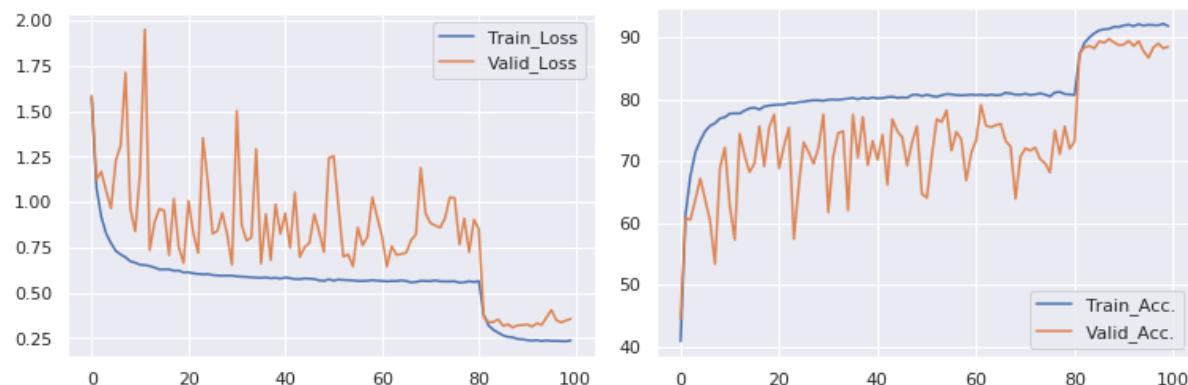
Learning Rate: 10 * Initial Learning Rate (1.e-4)

Epoch: [100 | 100] LR: 0.010000

Processing (391/391) Data: 0.020s | Batch: 0.062s | Total: 0:00:24 | ETA: 0:00:01 | Loss: 0.2369 | top1: 91.8020 | top5: 99.8460

Processing (100/100) Data: 0.010s | Batch: 0.028s | Total: 0:00:02 | ETA: 0:00:01 | Loss: 0.3560 | top1: 88.4500 | top5: 99.5700

Best acc: tensor (89.7200, device='cuda:0')



As we can see from the figures above, increasing the learning rate by 10 times will lead to increase the validation error which means we need enough epochs to fix this issue. And as a result, we got a lower validation accuracy 89.7200.

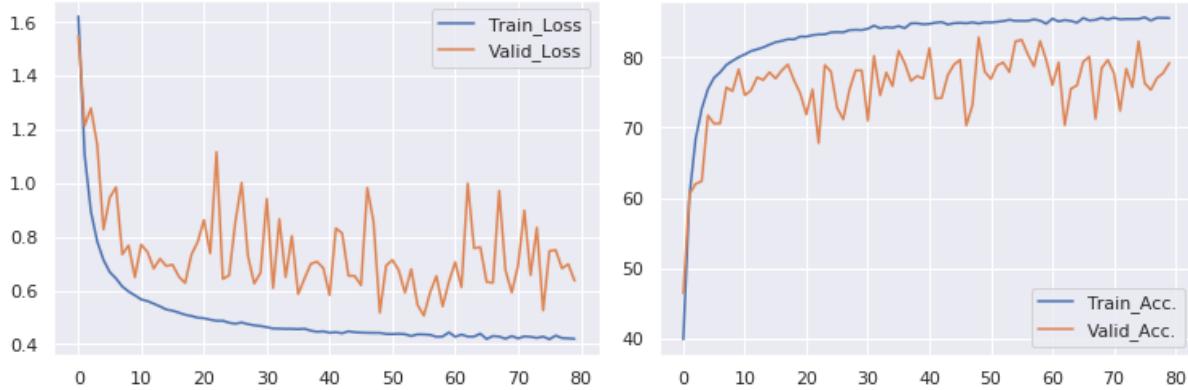
Learning Rate: 5 * Initial Learning Rate (1.e-4)

Epoch: [80 | 80] LR: 0.100000

Processing (391/391) Data: 0.018s | Batch: 0.063s | Total: 0:00:24 | ETA: 0:00:01 | Loss: 0.4200 | top1: 85.5440 | top5: 99.4540

Processing (100/100) Data: 0.009s | Batch: 0.028s | Total: 0:00:02 | ETA: 0:00:01 | Loss: 0.6364 | top1: 79.1600 | top5: 98.9100

Best acc: tensor (82.8200, device='cuda:0')



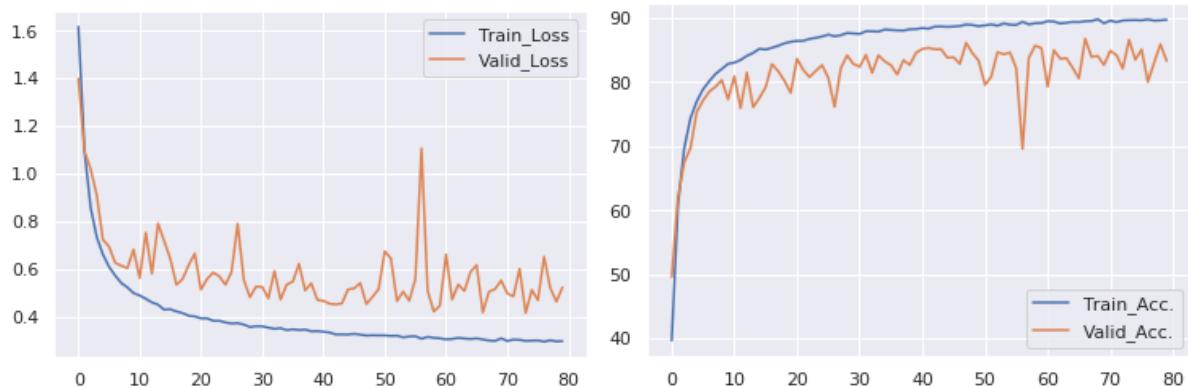
Learning Rate: 2 * Initial Learning Rate (1.e-4)

Epoch: [80 | 80] LR: 0.100000

Processing (391/391) Data: 0.019s | Batch: 0.063s | Total: 0:00:24 | ETA: 0:00:01 | Loss: 0.2966 | top1: 89.7520 | top5: 99.7340

Processing (100/100) Data: 0.009s | Batch: 0.029s | Total: 0:00:02 | ETA: 0:00:01 | Loss: 0.5221 | top1: 83.3100 | top5: 99.2800

Best acc: tensor (86.7800, device='cuda:0')



Learning Rate: 1/10 * Initial Learning Rate (1.e-4)

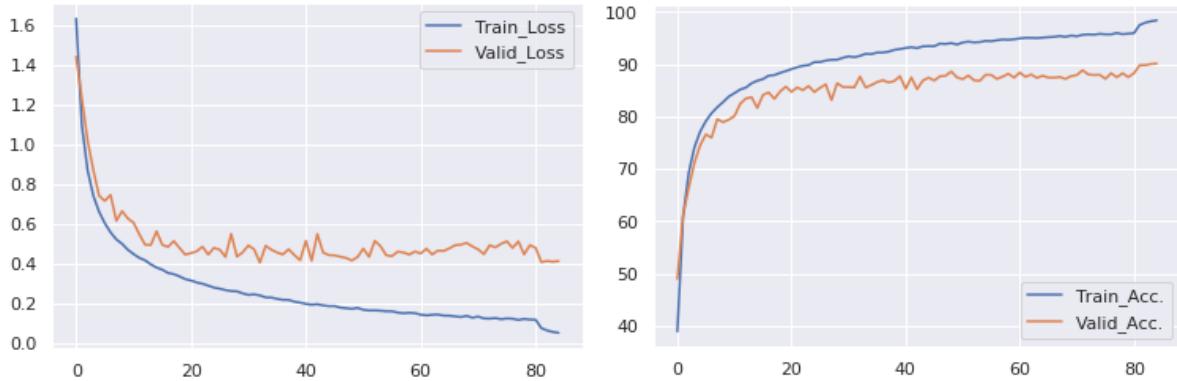
Epoch: [85 | 85] LR: 0.010000

Processing (391/391) Data: 0.020s | Batch: 0.063s | Total: 0:00:24 | ETA: 0:00:01 | Loss: 0.0488 | top1: 98.3680 | top5: 99.9980

Processing (100/100) Data: 0.009s | Batch: 0.030s | Total: 0:00:02 | ETA: 0:00:01 | Loss: 0.4103 | top1: 90.1200 | top5: 99.7200

Best acc: tensor (90.1200, device='cuda:0')

As seen in the figures below, decreasing the learning rate led to stable learning rate but the gap between training and a validation LOSS stays the same of worth.



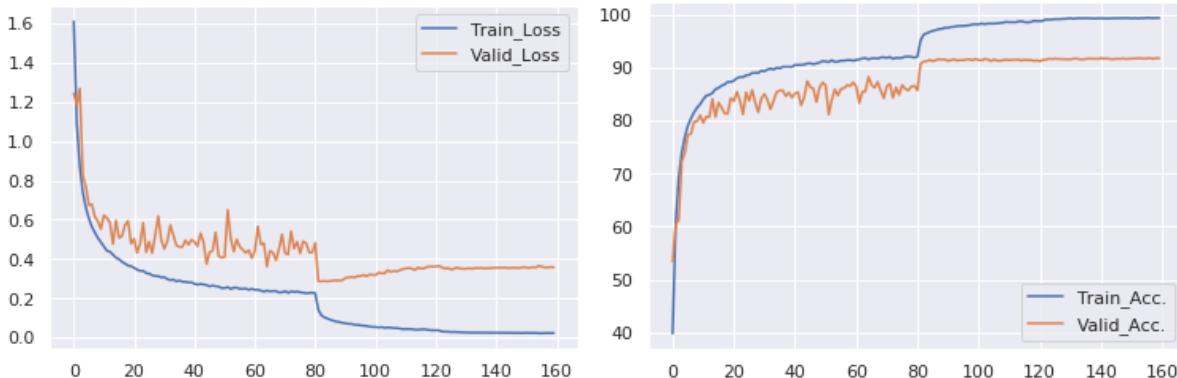
Increasing the number of the epochs to 160:

Epoch: [160 | 160] LR: 0.001000

Processing (391/391) Data: 0.021s | Batch: 0.063s | Total: 0:00:24 | ETA: 0:00:01 | Loss: 0.0228 | top1: 99.3800 | top5: 99.9980

Processing (100/100) Data: 0.008s | Batch: 0.029s | Total: 0:00:02 | ETA: 0:00:01 | Loss: 0.3576 | top1: 91.7600 | top5: 99.8000

Best acc: tensor (91.8200, device='cuda:0')



Increasing the number of epochs to a high value will lead to overfitting problem as it is obvious from the graph that the accuracy is not improving over more and more epochs

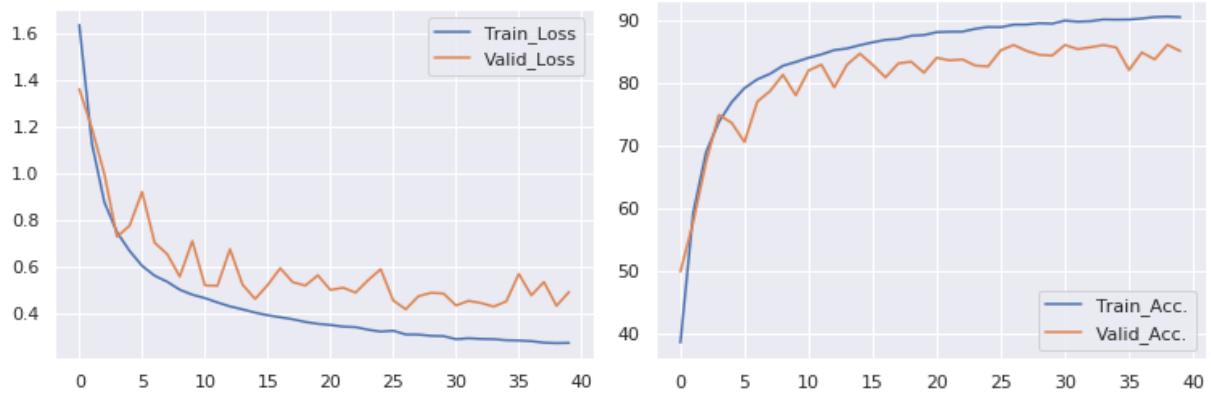
Decreasing the number of the epochs to 40:

Epoch: [40 | 40] LR: 0.100000

Processing (391/391) Data: 0.023s | Batch: 0.062s | Total: 0:00:24 | ETA: 0:00:01 | Loss: 0.2716 | top1: 90.4940 | top5: 99.8020

Processing (100/100) Data: 0.009s | Batch: 0.028s | Total: 0:00:02 | ETA: 0:00:01 | Loss: 0.4902 | top1: 85.0600 | top5: 99.2800

Best acc: tensor (86.0700, device='cuda:0')



Meanwhile, decreasing the number of epochs (in our case, 40) led to a bigger error range and a drop in the accuracy to 85%

ReadME DOC - Youshaa Murhij - Coding Homework 1

¬ Required Task 1

First I imported required libraries for printing and displaying results as Figures

```
1 %matplotlib inline  
  
1 import pandas as pd  
2 import matplotlib.pyplot as plt
```

Here I installed the lastest version of Pytorsh Platform

```
1 !pip install -q torch==1.3.0 torchvision  
  
1 %cd /content/drive/My\ Drive/  
  
1 import torch  
2 print(torch.__version__)  
  
⇒ 1.3.0+cu100
```

After that, we need to download the base line code for this paper from github

```
1 !git clone --recursive https://github.com/chenhuims/pytorch-classification.git  
  
1 %cd /content/drive/My\ Drive/pytorch-classification/
```

Here I called the train function on resnet and set the depth to 20 with 164 epochs and saved the re

```
1 %cd /content/drive/My\ Drive/pytorch-classification/  
2 !python cifar.py -a resnet --depth 20 --epochs 164 --schedule 81 122 --gamma 0.1 --wd 1
```

Required Libraries to display the graphs in a good way

```
1 import seaborn as sns  
2 sns.set()
```

Displaying the Loss and the Accuracy

```
1 %cd /content/drive/My\ Drive/pytorch-classification/checkpoint/cifar10/resnet-20/
```

```
1 %cd /content/drive/My Drive/pytorch-classification/checkpoint/cifar10/resnet-56/
2 file = pd.read_csv('log.txt',delimiter=' ')
3 plt.plot(file['Train_Loss'])
4 plt.plot(file['Valid_Loss'])
5 plt.legend(['Train_Loss','Valid_Loss'],loc='upper right');
6 plt.show()
```

Displaying the accuracy figure:

```
1 plt.plot(file['Train_Acc.'])
2 plt.plot(file['Valid_Acc.'])
3 plt.legend(['Train_Acc. ','Valid_Acc.'],loc='lower right');
4 plt.show()
```

Here I called the train function on resnet and set the depth to 56 with 125 epochs and saved the re

```
1 %cd /content/drive/My Drive/pytorch-classification/
2 !python cifar.py -a resnet --depth 56 --epochs 125 --schedule 81 122 --gamma 0.1 --wd 1
```

1 Displaying the Loss and the Accuracy

```
1 %cd /content/drive/My Drive/pytorch-classification/checkpoint/cifar10/resnet-56/
2 file = pd.read_csv('log.txt',delimiter=' ')
3 plt.plot(file['Train_Loss'])
4 plt.plot(file['Valid_Loss'])
5 plt.legend(['Train_Loss','Valid_Loss'],loc='upper right');
6 plt.show()
```

```
1 plt.plot(file['Train_Acc.'])
2 plt.plot(file['Valid_Acc.'])
3 plt.legend(['Train_Acc. ','Valid_Acc.'],loc='lower right');
4 plt.show()
```

Here I called the train function on CIFAR10 and set the depth to 110 with 110 epochs and saved th

```
1 %cd /content/drive/My Drive/pytorch-classification/
2 !python cifar.py -a resnet --depth 110 --epochs 110 --schedule 81 122 --gamma 0.1 --wd
```

Here I called the train function on cifar100 and set the depth to 20 with 100 epochs and saved the

```
1 %cd /content/drive/My Drive/pytorch-classification/
2 !python cifar.py -a resnet --dataset cifar100 --depth 20 --epochs 100 --schedule 81 122
```

Here I called the train function on cifar100 and set the depth to 56 with 100 epochs and saved the

```
1 %cd /content/drive/My Drive/pytorch-classification/
2 !python cifar.py -a resnet --dataset cifar100 --depth 56 --epochs 100 --schedule 81 122
```

Displaying the Loss and the Accuracy

```
1 %cd /content/drive/My Drive/pytorch-classification/checkpoint/cifar100/resnet-56/
2 file = pd.read_csv('log.txt',delimiter=' ')
3 plt.plot(file['Train_Loss'])
4 plt.plot(file['Valid_Loss'])
5 plt.legend(['Train_Loss','Valid_Loss'],loc='upper right');
6 plt.show()
```

```
1 plt.plot(file['Train_Acc.'])
2 plt.plot(file['Valid_Acc.'])
3 plt.legend(['Train_Acc. ','Valid_Acc.'],loc='lower right');
4 plt.show()
```

Here I called the train function on cifar100 and set the depth to 110 with 100 epochs and saved the log file.

```
1 %cd /content/drive/My Drive/pytorch-classification/
2 !python cifar.py -a resnet --dataset cifar100 --depth 110 --epochs 100 --schedule 81 122
```

▼ Optional Task 1

First, we multiply the base learning rate by 10 and check for changes in the results

```
1 %cd /content/drive/My Drive/pytorch-classification/
2 !python cifar.py -a resnet --dataset cifar10 --depth 20 --epochs 100 --schedule 81 122
```

Displaying the Loss and the Accuracy

```
1 %cd /content/drive/My Drive/pytorch-classification/checkpoint/cifar10_optional/resnet-2
2 file = pd.read_csv('log.txt',delimiter=' ')
3 plt.plot(file['Train_Loss'])
4 plt.plot(file['Valid_Loss'])
5 plt.legend(['Train_Loss','Valid_Loss'],loc='upper right');
6 plt.show()
```

```
1 plt.plot(file['Train_Acc.'])
2 plt.plot(file['Valid_Acc.'])
3 plt.legend(['Train_Acc. ','Valid_Acc.'],loc='lower right');
4 plt.show()
```

then, we multiply the base learning rate by 5 and check for changes in the results

```
1 %cd /content/drive/My Drive/pytorch-classification/
2 !python cifar.py -a resnet --dataset cifar10 --depth 20 --epochs 80 --schedule 81 122 -
```

Displaying the Loss and the Accuracy

```
1 %cd /content/drive/My Drive/pytorch-classification/checkpoint/cifar10_optional/resnet-2
2 file = pd.read_csv('log.txt',delimiter=' ')
3 plt.plot(file['Train_Loss'])
4 plt.plot(file['Valid_Loss'])
5 plt.legend(['Train_Loss','Valid_Loss'],loc='upper right');
6 plt.show()
```

```
1
```

```
1 plt.plot(file['Train_Acc.'])
2 plt.plot(file['Valid_Acc.'])
3 plt.legend(['Train_Acc.', 'Valid_Acc.'], loc='lower right');
4 plt.show()
```

we multiply the base learning rate by 2

```
1 %cd /content/drive/My Drive/pytorch-classification/
2 !python cifar.py -a resnet --dataset cifar10 --depth 20 --epochs 80 --schedule 81 122 -
```

we multiply the base learning rate by 1/10

```
1
```

```
1 %cd /content/drive/My Drive/pytorch-classification/
2 !python cifar.py -a resnet --dataset cifar10 --depth 20 --epochs 85 --schedule 81 122 -
```

Displaying the Loss and the Accuracy

```
1 %cd /content/drive/My Drive/pytorch-classification/checkpoint/cifar10_optional/resnet-2
2 file = pd.read_csv('log.txt',delimiter=' ')
3 plt.plot(file['Train_Loss'])
4 plt.plot(file['Valid_Loss'])
5 plt.legend(['Train_Loss','Valid_Loss'],loc='upper right');
6 plt.show()
```

```
1 plt.plot(file['Train_Acc.'])
2 plt.plot(file['Valid_Acc.'])
3 plt.legend(['Train_Acc.', 'Valid_Acc.'], loc='lower right');
4 plt.show()
```

we double the number of training epochs

```
1 %cd /content/drive/My Drive/pytorch-classification/
2 !python cifar.py -a resnet --dataset cifar10 --depth 20 --epochs 160 --schedule 81 122
```

we half the number of training epochs

```
1 %cd /content/drive/My Drive/pytorch-classification/
2 !python cifar.py -a resnet --dataset cifar10 --depth 20 --epochs 40 --schedule 81 122 -
```

Displaying the Loss and the Accuracy

```
1 %cd /content/drive/My Drive/pytorch-classification/checkpoint/cifar10_optional/resnet-2
2 file = pd.read_csv('log.txt',delimiter=' ')
3 plt.plot(file['Train_Loss'])
4 plt.plot(file['Valid_Loss'])
5 plt.legend(['Train_Loss','Valid_Loss'],loc='upper right');
6 plt.show()

1 plt.plot(file['Train_Acc.'])
2 plt.plot(file['Valid_Acc.'])
3 plt.legend(['Train_Acc.', 'Valid_Acc.'], loc='lower right');
4 plt.show()
```

▼ Training script for CIFAR-10/100

```
1
2
3 from __future__ import print_function
4
5 import argparse
6 import os
7 import shutil
8 import time
9 import random
10
11 import torch
12 import torch.nn as nn
13 import torch.nn.parallel
14 import torch.backends.cudnn as cudnn
15 import torch.optim as optim
16 import torch.utils.data as data
17 import torchvision.transforms as transforms
18 import torchvision.datasets as datasets
19 import models.cifar as models
20
21 from utils import Bar, Logger, AverageMeter, accuracy, mkdir_p, savefig
22
23
24 model_names = sorted(name for name in models.__dict__
25     if name.islower() and not name.startswith("__")
26     and callable(models.__dict__[name]))
27
28 parser = argparse.ArgumentParser(description='PyTorch CIFAR10/100 Training')
```

```

29 # Datasets
30 parser.add_argument('--d', '--dataset', default='cifar10', type=str)
31 parser.add_argument('--j', '--workers', default=4, type=int, metavar='N',
32                     help='number of data loading workers (default: 4)')
33 # Optimization options
34 parser.add_argument('--epochs', default=300, type=int, metavar='N',
35                     help='number of total epochs to run')
36 parser.add_argument('--start-epoch', default=0, type=int, metavar='N',
37                     help='manual epoch number (useful on restarts)')
38 parser.add_argument('--train-batch', default=128, type=int, metavar='N',
39                     help='train batchsize')
40 parser.add_argument('--test-batch', default=100, type=int, metavar='N',
41                     help='test batchsize')
42 parser.add_argument('--lr', '--learning-rate', default=0.1, type=float,
43                     metavar='LR', help='initial learning rate')
44 parser.add_argument('--drop', '--dropout', default=0, type=float,
45                     metavar='Dropout', help='Dropout ratio')
46 parser.add_argument('--schedule', type=int, nargs='+', default=[150, 225],
47                     help='Decrease learning rate at these epochs.')
48 parser.add_argument('--gamma', type=float, default=0.1, help='LR is multiplied by gamma')
49 parser.add_argument('--momentum', default=0.9, type=float, metavar='M',
50                     help='momentum')
51 parser.add_argument('--weight-decay', '--wd', default=5e-4, type=float,
52                     metavar='W', help='weight decay (default: 1e-4)')
53 # Checkpoints
54 parser.add_argument('--c', '--checkpoint', default='checkpoint', type=str, metavar='PATH'
55                     help='path to save checkpoint (default: checkpoint)')
56 parser.add_argument('--resume', default='', type=str, metavar='PATH',
57                     help='path to latest checkpoint (default: none)')
58 # Architecture
59 parser.add_argument('--arch', '-a', metavar='ARCH', default='resnet20',
60                     choices=model_names,
61                     help='model architecture: ' +
62                         ' | '.join(model_names) +
63                         ' (default: resnet18)')
64 parser.add_argument('--depth', type=int, default=29, help='Model depth.')
65 parser.add_argument('--block-name', type=str, default='BasicBlock',
66                     help='the building block for Resnet and Preresnet: BasicBlock, Bott'
67 parser.add_argument('--cardinality', type=int, default=8, help='Model cardinality (grou'
68 parser.add_argument('--widen-factor', type=int, default=4, help='Widen factor. 4 -> 64,'
69 parser.add_argument('--growthRate', type=int, default=12, help='Growth rate for DenseNe'
70 parser.add_argument('--compressionRate', type=int, default=2, help='Compression Rate (t'
71 # Miscs
72 parser.add_argument('--manualSeed', type=int, help='manual seed')
73 parser.add_argument('--e', '--evaluate', dest='evaluate', action='store_true',
74                     help='evaluate model on validation set')
75 #Device options
76 parser.add_argument('--gpu-id', default='0', type=str,
77                     help='id(s) for CUDA_VISIBLE_DEVICES')
78
79 args = parser.parse_args()
80 state = {k: v for k, v in args._get_kwargs()}
81
82 # Validate dataset
83 assert args.dataset == 'cifar10' or args.dataset == 'cifar100', 'Dataset can only be ci

```

```

84
85 # Use CUDA
86 os.environ['CUDA_VISIBLE_DEVICES'] = args.gpu_id
87 use_cuda = torch.cuda.is_available()
88
89 # Random seed
90 if args.manualSeed is None:
91     args.manualSeed = random.randint(1, 10000)
92 random.seed(args.manualSeed)
93 torch.manual_seed(args.manualSeed)
94 if use_cuda:
95     torch.cuda.manual_seed_all(args.manualSeed)
96
97 best_acc = 0 # best test accuracy
98
99 def main():
100     global best_acc
101     start_epoch = args.start_epoch # start from epoch 0 or last checkpoint epoch
102
103     if not os.path.isdir(args.checkpoint):
104         mkdir_p(args.checkpoint)
105
106
107
108     # Data
109     print('==> Preparing dataset %s' % args.dataset)
110     transform_train = transforms.Compose([
111         transforms.RandomCrop(32, padding=4),
112         transforms.RandomHorizontalFlip(),
113         transforms.ToTensor(),
114         transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
115     ])
116
117     transform_test = transforms.Compose([
118         transforms.ToTensor(),
119         transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
120     ])
121     if args.dataset == 'cifar10':
122         dataloader = datasets.CIFAR10
123         num_classes = 10
124     else:
125         dataloader = datasets.CIFAR100
126         num_classes = 100
127
128
129     trainset = dataloader(root='./data', train=True, download=True, transform=transform_train)
130     trainloader = data.DataLoader(trainset, batch_size=args.train_batch, shuffle=True,
131
132     testset = dataloader(root='./data', train=False, download=False, transform=transform_test)
133     testloader = data.DataLoader(testset, batch_size=args.test_batch, shuffle=False, nu
134
135     # Model
136     print("==> creating model '{}'".format(args.arch))
137     if args.arch.startswith('resnext'):
138         model = models.__dict__[args.arch](

```

```

139             cardinality=args.cardinality,
140             num_classes=num_classes,
141             depth=args.depth,
142             widen_factor=args.widen_factor,
143             dropRate=args.drop,
144         )
145     elif args.arch.startswith('densenet'):
146         model = models.__dict__[args.arch](
147             num_classes=num_classes,
148             depth=args.depth,
149             growthRate=args.growthRate,
150             compressionRate=args.compressionRate,
151             dropRate=args.drop,
152         )
153     elif args.arch.startswith('wrn'):
154         model = models.__dict__[args.arch](
155             num_classes=num_classes,
156             depth=args.depth,
157             widen_factor=args.widen_factor,
158             dropRate=args.drop,
159         )
160     elif args.arch.endswith('resnet'):
161         model = models.__dict__[args.arch](
162             num_classes=num_classes,
163             depth=args.depth,
164             block_name=args.block_name,
165         )
166     else:
167         model = models.__dict__[args.arch](num_classes=num_classes)
168
169 model = torch.nn.DataParallel(model).cuda()
170 cudnn.benchmark = True
171 print('    Total params: %.2fM' % (sum(p.numel() for p in model.parameters())/10000
172 criterion = nn.CrossEntropyLoss()
173 optimizer = optim.SGD(model.parameters(), lr=args.lr, momentum=args.momentum, weigh
174
175 # Resume
176 title = 'cifar-10-' + args.arch
177 if args.resume:
178     # Load checkpoint.
179     print('==> Resuming from checkpoint..')
180     assert os.path.isfile(args.resume), 'Error: no checkpoint directory found!'
181     args.checkpoint = os.path.dirname(args.resume)
182     checkpoint = torch.load(args.resume)
183     best_acc = checkpoint['best_acc']
184     start_epoch = checkpoint['epoch']
185     model.load_state_dict(checkpoint['state_dict'])
186     optimizer.load_state_dict(checkpoint['optimizer'])
187     logger = Logger(os.path.join(args.checkpoint, 'log.txt'), title=title, resume=T
188 else:
189     logger = Logger(os.path.join(args.checkpoint, 'log.txt'), title=title)
190     logger.set_names(['Learning_Rate', 'Train_Loss', 'Valid_Loss', 'Train_Acc.', 'V
191
192 if args.evaluate:

```

```

194     print('\nEvaluation only')
195     test_loss, test_acc = test(testloader, model, criterion, start_epoch, use_cuda)
196     print(' Test Loss: %.4f, Test Acc: %.2f' % (test_loss, test_acc))
197     return
198
199 # Train and val
200 for epoch in range(start_epoch, args.epochs):
201     adjust_learning_rate(optimizer, epoch)
202
203     print('\nEpoch: [%d | %d] LR: %f' % (epoch + 1, args.epochs, state['lr']))
204
205     train_loss, train_acc = train(trainloader, model, criterion, optimizer, epoch,
206     test_loss, test_acc = test(testloader, model, criterion, epoch, use_cuda)
207
208     # append logger file
209     logger.append([state['lr'], train_loss, test_loss, train_acc, test_acc])
210
211     # save model
212     is_best = test_acc > best_acc
213     best_acc = max(test_acc, best_acc)
214     save_checkpoint({
215         'epoch': epoch + 1,
216         'state_dict': model.state_dict(),
217         'acc': test_acc,
218         'best_acc': best_acc,
219         'optimizer': optimizer.state_dict(),
220     }, is_best, checkpoint=args.checkpoint)
221
222 logger.close()
223 logger.plot()
224 savefig(os.path.join(args.checkpoint, 'log.eps'))
225
226 print('Best acc:')
227 print(best_acc)
228
229 def train(trainloader, model, criterion, optimizer, epoch, use_cuda):
230     # switch to train mode
231     model.train()
232
233     batch_time = AverageMeter()
234     data_time = AverageMeter()
235     losses = AverageMeter()
236     top1 = AverageMeter()
237     top5 = AverageMeter()
238     end = time.time()
239
240     bar = Bar('Processing', max=len(trainloader))
241     for batch_idx, (inputs, targets) in enumerate(trainloader):
242         # measure data loading time
243         data_time.update(time.time() - end)
244
245         if use_cuda:
246             inputs, targets = inputs.cuda(), targets.cuda(async=True)
247             inputs, targets = torch.autograd.Variable(inputs), torch.autograd.Variable(targ
248

```

```

249     # compute output
250     outputs = model(inputs)
251     loss = criterion(outputs, targets)
252
253     # measure accuracy and record loss
254     prec1, prec5 = accuracy(outputs.data, targets.data, topk=(1, 5))
255     if float(torch.__version__[:3]) < 0.5:
256         losses.update(loss.data[0], inputs.size(0))
257         top1.update(prec1[0], inputs.size(0))
258         top5.update(prec5[0], inputs.size(0))
259     else:
260         losses.update(loss.data, inputs.size(0))
261         top1.update(prec1, inputs.size(0))
262         top5.update(prec5, inputs.size(0))
263
264     # compute gradient and do SGD step
265     optimizer.zero_grad()
266     loss.backward()
267     optimizer.step()
268
269     # measure elapsed time
270     batch_time.update(time.time() - end)
271     end = time.time()
272
273     # plot progress
274     bar.suffix = '({batch}/{size}) Data: {data:.3f}s | Batch: {bt:.3f}s | Total: {'
275             batch=batch_idx + 1,
276             size=len(trainloader),
277             data=data_time.avg,
278             bt=batch_time.avg,
279             total=bar.elapsed_td,
280             eta=bar.eta_td,
281             loss=losses.avg,
282             top1=top1.avg,
283             top5=top5.avg,
284             )
285     bar.next()
286 bar.finish()
287 return (losses.avg, top1.avg)
288
289 def test(testloader, model, criterion, epoch, use_cuda):
290     global best_acc
291
292     batch_time = AverageMeter()
293     data_time = AverageMeter()
294     losses = AverageMeter()
295     top1 = AverageMeter()
296     top5 = AverageMeter()
297
298     # switch to evaluate mode
299     model.eval()
300
301     end = time.time()
302     bar = Bar('Processing', max=len(testloader))
303     for batch_idx, (inputs, targets) in enumerate(testloader):
304         " "

```

```

304     # measure data loading time
305     data_time.update(time.time() - end)
306
307     if use_cuda:
308         inputs, targets = inputs.cuda(), targets.cuda()
309         inputs, targets = torch.autograd.Variable(inputs, volatile=True), torch.autogra
310
311     # compute output
312     outputs = model(inputs)
313     loss = criterion(outputs, targets)
314
315     # measure accuracy and record loss
316     prec1, prec5 = accuracy(outputs.data, targets.data, topk=(1, 5))
317     if float(torch.__version__[:3]) < 0.5:
318         losses.update(loss.data[0], inputs.size(0))
319         top1.update(prec1[0], inputs.size(0))
320         top5.update(prec5[0], inputs.size(0))
321     else:
322         losses.update(loss.data, inputs.size(0))
323         top1.update(prec1, inputs.size(0))
324         top5.update(prec5, inputs.size(0))
325
326     # measure elapsed time
327     batch_time.update(time.time() - end)
328     end = time.time()
329
330     # plot progress
331     bar.suffix = '({batch}/{size}) Data: {data:.3f}s | Batch: {bt:.3f}s | Total: {'
332                 batch=batch_idx + 1,
333                 size=len(testloader),
334                 data=data_time.avg,
335                 bt=batch_time.avg,
336                 total=bar.elapsed_td,
337                 eta=bar.eta_td,
338                 loss=losses.avg,
339                 top1=top1.avg,
340                 top5=top5.avg,
341                 )
342     bar.next()
343 bar.finish()
344 return (losses.avg, top1.avg)
345
346 def save_checkpoint(state, is_best, checkpoint='checkpoint', filename='checkpoint.pth.t
347     filepath = os.path.join(checkpoint, filename)
348     torch.save(state, filepath)
349     if is_best:
350         shutil.copyfile(filepath, os.path.join(checkpoint, 'model_best.pth.tar'))
351
352 def adjust_learning_rate(optimizer, epoch):
353     global state
354     if epoch in args.schedule:
355         state['lr'] *= args.gamma
356         for param_group in optimizer.param_groups:
357             param_group['lr'] = state['lr']
358
359 if __name__ == '__main__':
360     main()

```

```
360     main()
```